

A Distributed, Collaborative Development Project of
a Policy Engine for Use in Buildings
Global Software Development

Berntsen, R., `raber@itu.dk`
Hansen, K., `kben@itu.dk`
Kokholm, T., `tkok@itu.dk`
Stanciulescu, S., `scas@itu.dk`
Wainach, N., `nicl@itu.dk`

May 14, 2013

Abstract

This paper presents a distributed, collaborative software prototype development project in the domain of building automation. The prototype system — the policy engine – is used for regulating the internal environment of a building. The global team behind the prototype consists of members from Strathmore University, Kenya and the IT-University of Copenhagen, Denmark.

The field of building automation has seen a lot of growth in recent years. Building automation can save energy and resources by e.g. turning off light or regulating the temperature when nobody is in the room. Conservation of energy and natural resources through building automation offers vast improvements compared to buildings where this automation is not deployed.

The approach to the project is to analyze, design and implement the policy engine prototype in a distributed, collaborative project. First we created the communication platform for the project. Then we managed the project while simultaneously developing the prototype. During the length of the project we founded multiple subgroups in regards to team members preferences and interests - seeking to optimize both our quality and quantity of code.

The end result was a fully working prototype of a policy engine platform, offering functionality to create, edit and execute policies. The engine was validated through a series of tests, including low level functionality tests — unit tests — and high level user tests.

We argue that our approach towards the development of the prototype is both valid and appropriate in relation to the goals, requirements and audience.

Contents

1	Introduction	1
1.1	Context	3
1.2	Problem	4
1.3	Learning Goals	4
1.4	Requirements	5
2	Related Work	6
3	Collaboration Strategy	8
3.1	Workflow	8
3.1.1	Stages	9
3.1.1.1	Initial Stage	9
3.1.1.2	Development Stage	9
3.1.1.3	Deployment Stage	10
3.1.2	Iterations	10
3.1.2.1	First Iteration	10
3.1.2.2	Second Iteration	11
3.1.2.3	Third Iteration	11
3.2	Collaboration	12
3.2.1	Create a common platform	12

3.2.2	Global Collaboration Gone Wrong	13
3.2.3	Tools	14
3.2.4	Mitigation	14
3.3	Technology	14
4	Design	17
4.1	Architecture	18
4.1.1	The management website	20
4.1.2	The policy engine	20
4.1.3	The database	22
4.1.4	The building simulator	22
5	Implementation	23
5.1	System	23
5.1.1	Interfaces	25
5.1.2	Domain	25
5.2	Persistence	26
5.3	Building querying	26
5.4	Executing policies	27
5.5	Front-end	28
5.5.1	Initial Ideas	28
5.5.2	User Interface	30
5.5.2.1	Intuitive Design	30
5.5.2.2	Flow	31
5.5.2.3	Response Times	32
5.5.2.4	Color Coding	32
5.5.2.5	Responsive Design	33

5.5.3	Policy Editor	33
5.5.3.1	Managing Policies	33
5.5.4	Use Case View	34
5.5.5	Usability	35
5.5.5.1	Browser Compatibility	35
5.5.5.2	Layers of Abstraction	36
6	Evaluation	37
6.1	Policy Engine System Evaluation	37
6.1.1	Log Testing	37
6.1.2	JUnit Testing	38
6.2	Usability Test	38
6.2.1	Think Aloud Protocol	38
6.2.1.1	Tasks	39
6.2.1.2	Results and Comments	41
6.2.1.3	Subset	42
7	Discussion	43
7.1	Collaboration	43
7.2	Prototype system	44
7.2.1	Future Work & Improvements	45
7.3	Project	45
8	Conclusion	48
	References	52
A	Appendix	53
A.1	Collaboration	53

A.1.1	Social Context	53
A.1.1.1	Common ground	53
A.1.1.2	Trust and First Impression	54
A.1.1.3	Collaboration Readiness	54
A.1.1.4	Ethnocentrism	55
A.1.1.5	Coupling of work	56
A.1.2	Collaborative work	56
A.1.2.1	Articulation work	56
A.1.2.2	Coordination	57
A.1.3	Groupware Technologies	57
A.1.3.1	Tools	58
A.1.3.2	Critical mass	58
A.1.3.3	Adaptation & Adaption Process	58
A.1.4	Virtual Project Management	59
A.1.4.1	Discontinuities	59
A.1.4.2	Continuities	60
A.2	Requirements	60
A.3	Tests	60
A.4	Project Plan	60
A.5	Project Plan	60

Chapter 1

Introduction

The research field of building- and home automation experiences a lot of growth – not least because of the promise to reduce energy consumption by more intelligent control, but also due to the possibility of heightened human comfort. Constructing resource efficient buildings makes sense, both in a political and economical perspective. In [18] it is stated that residential buildings use about 82% in total energy consumption on space and water heating. Electric appliances uses 11%. One will be able to save on these sources by using policies. Policies can be seen as a general rule that determines what to do in sense of controlling lights, heat, blinds and the likes. An example of a policy could be to turn off heating at midnight. Manually controlling an energy-saving policy is a time-consuming and inefficient task and the user have to know a lot of different equipment as well as information about the building. It is, however, achievable using building automation.

Regulating the environment of buildings to heighten human comfort, requires that building components can communicate and be controlled in a fashion that seems intuitive for human beings. Buildings today might come equipped with a suite of sensors and actuators, opening up for a degree of customizable control. Our collective need is that buildings can adapt to the users and the sensor-perceived environment. We use the term policy to describe a building automation program, and the term policy engine as the overall software entity that controls the aforementioned policies.

Trivial policies can be based on semi-static data, like time and weekdays. However this can have unforeseen and unwanted consequences. For example, a policy governing lightning activated merely by a static time schedule, might entail problems for people attending a rarely occurring late-night party in the building. If the event calendar of the building is accessible to the policy engine, a conditional

event-checking statement might ensure continuous lighting. However, in order to achieve a more fine grained control, sensory input is needed. The interaction with these policies are thought of and developed in respect of a facility management¹. The task of controlling the policies are also defined as being part of the tasks of FM.

By employing a policy engine, with access to the buildings sensors and actuators, both the building owner, the users and the administrators of the building benefits from the automation provided. If policies are correctly defined, building owners save energy and other natural resources while providing extra comfort to their tenants. A buildings users can experience a building autonomously adjusting its internal environment to suit their comfort and needs, while FM can achieve fine-grained control of the building.

As stated in [7], efficient energy saving is only possibly with building automation, known as “BA”: *Worldwide, there is no doubt that efficient energy saving is only possible with modern BA based on networking in all levels of abstraction.*

This project was defined in the course Global Software Development at IT University of Copenhagen². Our global development team consists of 5 students from ITU and 4 from Strathmore University in Nairobi, Kenya. We have been provided a Building Simulator, making up for the lack of a real building. The complexity with integrating to hardware and communication protocols are thus avoided, however the simulator has provided other complexities as will be evident later in this paper.

The development focus of this project is therefore geared towards this simulator. The end product is a web-based management application, that allows for centralized flow control between sensors and actuators in the simulated building. This is achieved by implementing a policy engine that allows for automated actuator responses based on sensor feedback. An example of this would be closing the blinds in excess sunlight or turning off the heater when the windows are open.

In this paper we will:

1. Distil requirements from course provided material.
2. Literature (re-)search on policy engines.
3. Develop a software solution that implements these requirements.
4. Document the collaborative project between the IT University of Copenhagen, Denmark and Strathmore University, Nairobi Kenya.

¹from now on FM

²From now on ITU

1.1 Context

Modern buildings get more and more complex – from the type of materials being used, to the services and infrastructure they provide. Our focus in this paper is building automation through the use of governing policies, and therefore economically speaking, relates to the buildings running costs, without lack of focus on human comfort. Today the cost of resources such as gas, diesel and electricity have been generally climbing for the last many years, making it still more important to use them with care. Another cost related to the usage of energy and natural resources are the taxes which also are on the rise. However, economy is not the only factor for saving the planets resources. Today it is considered so good PR to 'go green' that some companies build a virtual presence for this subject alone [13] [10] [32].

A part of the whole 'go green' concept is to control your energy and natural resource usage. Many companies have buildings that are either heated or cooled (or both), have lighting, appliances and server rooms to name a few. They may already be equipped intelligent heating systems, HVAC systems³ and AC's⁴ but many of these systems are typically proprietary and impossible—or very hard—to integrate with. As mentioned earlier this project integrates into a Building Simulator. This makes our task more simple compared to real work applications. We do not have to worry about actual sensor- and actuator hardware, wiring, costs, communication protocols, existing protocols, coverage of wireless communication signals and so forth.

We assume that the Building Simulator *is* the building. Our system is therefore *open* compared to proprietary, since all it takes to access, e.g. the temperature in a room is a http get operation with the building id and sensor in question. A everyday example could be a janitor, who might need to go to a room physically to check its temperature, and adjust the heater in that room, several times a day to achieve and maintain the desired temperature. Our governing policies makes it up for a building janitor, or at least the trivial work with adjusting room temperatures and the likes. Human staff starts out by defining the policies that the building needs. Then we can imagine the virtual janitor, issuing http get operations all day, to check the sensors mentioned in those policies. The policies consists of sensor-conditional behaviour that will adjust the actuators continuously. This is done by issuing http post commands to the building simulator with sensor id and the desired value.

³heating, ventilation, and air conditioning

⁴Air Conditioning

1.2 Problem

With the ability to manage every sensor in a given building, one can assume that the task of controlling these gets progressively—for each number of additional sensor and actuator—more complex. Any additional concern will add further complexity, such as optimizing the control to prioritize both economic and human comfort. One needs to consider this challenge: What good is a highly advanced building if the task of controlling it is just as complex?

We seek to design a system where FM can define a set of policies they believe is needed, based on their experience with the building. We do not expect FM to be IT experts, but we do expect them to be able to use a modern browser and be familiar with the sensor id's used in the building.

Our main consideration in this paper is: *Providing a web-based infrastructure to visually define, manage and monitor the scheduling of governing policies for a building. The policies should be executed at regular intervals, based on their time specific settings. The infrastructure should be provided in an easy accessible interface that provides FM with the means of controlling the building's actuator systems, based on sensory input and conditions expressed in the policies. It should be possible to define policies that will maintain human comfort in that building, while conserving energy with regards to heating and/or lighting.*

In addition to the general problem from above, we further want to point out these challenges:

1. The rating of proposed solution's usability.
2. Policy time control and manual override.
3. The use of complex policy data structures on both frontend and backend.

1.3 Learning Goals

The project in itself consists of two widely different aspects. They are the global software development and the policy engine. The new aspect in the project life cycle is for us how to collaborate with team members distributed in different continents and countries, with many differences and thus challenges to overcome.

These two aspects are also the ones the main learning goals are derived from. In this regard the goals can be grouped into two different parts. The first being how to collaborate across different countries and the challenges that this potentially could cause, including how to manage the team, how to communicate, how to

increase performance and the likes. The secondary being how we can solve the technical project and deliver a working prototype of a policy engine, including creating the domain model, data structure, creating the actual policy engine in a object oriented-manner and tying this together with the visual and user friendly frontend.

We could summarize the learning goals to be:

- How to work together in a distributed team and reflect on this.
- How to develop a policy engine that conforms to the stated requirements (See section 1.4).

1.4 Requirements

The requirements of the project are stated in a very open-ended manner, with only a few descriptive requirements to the product and process. These are formulated as requirements to the product and process. They are:

Product

- The students analyse their solution from their chosen sensors and actuators requirements (and additional requirements they can think of) into functionality for their application (30%).
- The students must design and implement a Web/Android application to monitor, control and visualize policies in a building (20%).
- The students must describe and evaluate their solution as used by facility people as metrics. Additional metrics can be considered and will be taking into consideration by the examiners (20%).

Process

- The students should be able to write a proper, understandable and well organized report (10%).
- The students should be able to reflect on a real world collaborative experience (20%).

Chapter 2

Related Work

A lot of research has emerged lately in the field of energy management systems for smart buildings. A similar work with the one presented in this paper has been done by Tiberkak et. al [36], where a Policy Based Architecture for Home Automation System is developed. The system is composed of the following sub-systems: one responsible for home automation, one for the local management of rooms, one for storing data, and a system for enabling communication between the first two sub systems. Different profiles are defined to improve energy efficiency. The concept of preferred and required profile is introduced, to differentiate between the preferences of a inhabitant and the maximum and minimum values of each environmental parameter in the required profile. Notifications and messages are sent between the layers when there is a change in the status of a room, and a appropriate decision is taken. Their approach is focused more on the confort of a person in the house than our approach, which is focused more on the management of a bigger building with more sensors and actuators. This approach could be used if a room in a building is known to be used frequently by the same persons, so one could deffine its own profile or policy. Their notification system is too complicated to use and for our needs, so we have employed a simple way to execute policies which does not need a communication protocol to be defined.

Another approach is the one taken by Li et. al [24] where they implement a energy management system for homes, that provides automatic metering and capability of taking decisions based on monitoring energy consumption. Tasks can be used to specify different actions required at different moments. A simulation has been done for 1000 homes and by using their system, they achieved a significant energy reduction.

In [16] a complete system for home management is described. Using ZigBee

and IEEE 802.15.4, it is easy to add new devices, connect them with already existing ones and control them using a remote device. However, no kind of automatic management exists for the light service, heating and air conditioning. This research would be useful for our work if there would be the need to add more sensors or actuators to the building. Our system should identify automatically the new installed sensors and actuators.

In [39] develop a prototyping platform based on Building Controls Virtual Test Bed framework [35] for controlling and testing networked sensors and actuators on physical system. An algorithm for controlling the light and blind system in the room has been developed. The system is configured to provide an illumination of 500 lux between 6 AM and 6 PM. This has been achieved by measuring the daylight and setting the blind to block direct sun beam into the room. The system managed to achieve a reduction by an average of 17% of cooling demands and a maximum of 57% of lighting demand compared to the reference room. This has been optimized to provide the optimum illumination in a room

Krioukov et al [21] take another approach with their Building Stack Application. They provide an application programming interface and runtime for portable building applications. Developers can express their intent in a easy way e.g. "turn off the lights for top floor cubicles near windows" or can handle building differences by having support for programatic exploration of the building's components. It consists of three layers, the query based API layer, the driver layer and the abstract interconnection layer. The query layer provides functionality for selecting objects based on attributes, type and functional or spatial relationships. The driver layer is composed of high-level and low-level drivers to encompass the number of different components in a given building. The low-level defines the types of the components and the high-level provides descriptive information of those components. The asbstract interconnection layer is a RESTful interface to facilitate communication between the application's request and the building control protocols.

Chapter 3

Collaboration Strategy

In this chapter we present our tactics and methods in regards of the collaboration. The level of abstraction is primarily on a descriptive level, while some challenges are also reflected upon.

3.1 Workflow

In our design of a policy engine, we decided to use an iterative design approach, both for the frontend and backend. Using this approach we were able to continuously evaluate both the concept and our software. We have used informal testing throughout the complete project. The goal was to continuously test and optimize the software, learn about the stability and performance, and finally to adjust the implementation according to these findings. An iterative design approach is commonly used for the frontend design. This allows designers to identify any usability issues that may arise in the user interface before it is put into wide use. Even the best usability experts cannot design perfect user interfaces in a single attempt, so a usability engineering life cycle should be built around the concept of iteration [26].

The figure below (fig. 3.1) illustrates the different iterations, phases and stages of our project.

Each of these illustrated elements were important for the process and contains all valuable information. The figure contains a timeline, which chronologically shows how the project evolved. The project went through three overall stages. These are known as the “Initial Stage”, “Development Stage” and “Closing Stage”.




Figure 3.1: A generalized view of our workflow process

3.1.1 Stages

3.1.1.1 Initial Stage

Before kicking off the project and delve into the actual development, we needed to figure out how to handle the group as *one* collaborating team. This part was of high importance due to the fact that we, as students from two different continents, did not know each other at all. The task was complicated by this. One could presume that in a normal work setting at least the different roles in the team would already have been settled. Some are hired to do the backend (developer) and others others to manage the project (project manager). We needed a specific initial stage to cover this part of getting to know each other, figure out roles and knowledge areas. This was internally known as the “Initial Stage”. The stage created the necessary basis to work from and we could start to move the project in the right direction.

3.1.1.2 Development Stage

During the development stage several iterations of analysis, design and implementation were executed. As the figure shows, the repetitive phases in the development stage is an iterative process. Even though we were through this loop multiple times, we will only describe the three major iterations, in the section below (See sec. 3.1.2). These iterations are coherent with the project plan and present our general approach to the project. All of the three iterations contains

multiple inclusive —by assessing the artefact—simpler iterations. Describing each iteration in detail would be a somewhat cumbersome process.

The reasoning for choosing three major iterations was to accommodate the overall deadlines for the project. Each deadline requested different material containing information about different levels of the project. The first deadline was focused on the requirements and motivation. The second delivery focused on related work and model. Third delivery's focus was on the actual product and its code. These deadlines were a part of our overall project plan and thus led to having three overall deadlines. Initially our goal was to have four major iterations but this was changed due to our unsuccessful efforts to include our colleagues from Kenya. The postponement of deadlines and a general patient attitude caused us to eliminate one major iteration.

The development stage contains furthermore a group coordination phase. The coordination is reached multiple times throughout the iterations. During this phase important topics were covered, such as sharing status, separation and distribution of tasks, figure out how to proceed, discuss upcoming challenges and the likes. Multiple collaboration methods and tools were used throughout this coordination phase (See 3.2 and 3.2.3).

3.1.1.3 Deployment Stage

The deployment stage was considered the simplest and easiest to handle stage of our process. Upon reaching agreement that the development was finalized, the product was deployed to our server. This concluded our work on the product and finalized the development.

3.1.2 Iterations

3.1.2.1 First Iteration

Our focus during the first iteration was to get the group work to flow and figure out how to handle the data services provided by the Danish university. Most of the iteration was related to get to know each other and figure out how to move the project forward. The artefacts produced here was of lower level, such as diagrams of the group members programming skills, use case diagrams, first thoughts of domain model, early prototypes of the frontend interface and early implementations of the data services.

It was difficult to really push the project forward at this phase. We used different means of communication with the Kenyans, but without any positive result. Our

best result was reached when two out of four African members replied to our emails. No feedback, replies, comments or the likes was received from the two last members.

3.1.2.2 Second Iteration

During the second major phase the groups communication platform and general condition was stable - but unfortunately without any involvement from our Kenyan team members. We knew that we had to push the project forward, and had already wasted too much time on the effort to engage and activate unresponsive members. We were behind our project plan in an attempt to give them enough time to respond, without any luck. We tried to activate one semi-active member from Kenya as a team leader, but his response was, that he was not able to contact the other members.

Our focus changed from here. At the following team meeting we agreed upon, that it would be up to us, to get the project going. We chose to do so because of the fact that we did not receive any kind of feedback from Kenya. We would still let them know every step in the project, and we would continue to try to involve them as much as possible, by emailing them and the likes. We developed multiple artefacts at this point. Our goal through the second iteration was to get a working backend up and running. To do this, we had to choose the best fitting development platform and hereby programming language. The group had to agree on a set of requirements for the engine derived from the requirements in the assignment. The data model was completed and implemented, development of an early working prototype of the backend, and generated several jUnit tests to validate its robustness. Furthermore the work on the frontend started to take form. We did our first evaluation of the frontend (see 6.2).

3.1.2.3 Third Iteration

The third and final iteration was a hectic period with focus on getting things done and complete tasks. Because of the earlier effort to activate the Kenyan members, and the time consumed here, we were lacking behind. The now only moderate active member from Kenya chose to leave the group due to personal circumstances. The frontend was implemented with support from data provided by the backend. A second usability test was produced. Also this paper was mostly produced here, by combining bits and pieces and adding lots of new material. Finally the end product was deployed to our server.

During all of the iterations, both major and minor, we used a palette of different methods and tools. These are highly linked with the project phases. The most

important are introduced and explained by usage, in the parts that follow below.

3.2 Collaboration

A set of different methods were used during the initial stage, in the group establishment, and afterwards in the development stage. These were applied in order to promote and streamline collaboration. Several challenges were met and could have been solved by applying the correct mitigation strategy. The collaboration between the Danish and Kenyan team never evolved enough to engage some of the tactics. The following sections will only highlight the actual engaged methods, and will leave out some methods that are only used briefly. This is done to inform which important actions were taken to engage the Kenyan members and not to blur the picture by which methods that could have been used at a later stage, if the collaboration evolved. For a complete reference of all methods that we sought to use and how they could have been used, see Appendix A.1.

3.2.1 Create a common platform

At the time of the project start we were fully aware that we had several discontinuities, as described by Watson-Manheim et al. [38], between the two groups. Too mention a few, the most challenging were, that we did not share the same culture, time zone —Kenya was ahead of two hours—, possibly skills, history and the likes. All of which makes it a challenge to collaborate. One way to solve this is to get to know the different team members in the group. The first step we took in order to make the group work as one unit was to try to create a common ground. The term to create common ground “refers to that knowledge that the participants have in common, and they are aware that they have it in common” [29]. The fact that the members in the teams, local as well as global, had not met each other, or collaborated, before, only stressed the importance of this. In a group like this, where members are spread across different continents, the common ground and the social context that ones behave in, is important. No one knew what to expect from each other, and different cultures introduces different challenges. Our first approach to mitigate this problem was to email each member, by the email supplied by the coordinators, and introduce the members from ITU. The introduction was superficial and not exhaustive, in order not to frighten anybody. Basic information was listed, such as name, gender, age, contact methods and the likes. This would be a 30 seconds email to compose, and we nonetheless never received any replies. To mitigate this challenge of no replies, we tried to create a spreadsheet listing these simple informations, and

expanded it with individual skills significant to the project. It took the Kenyan members close to two weeks to complete this form.

The fact that some Kenyan members took several days, and weeks, to reply our emails created a tense feeling in the group. The Danish group had high hopes for the collaboration and were left with a perception of sluggishness towards the Kenyan group. This caused that several members lost trust in the African members, as introduced by Jarvenpaa et al. [19]. This shows how important the first impression is.

To make it easy to communicate between the group, we created an email list, that every member was a part of. This was done to enable members to create a global awareness [23] in the team and inform each other which challenges we currently faced. We informed every member to use this email. Only a few mails from one member in Kenya was received. The replies were often late and inadequate.

3.2.2 Global Collaboration Gone Wrong

Even though we were only in the early stages of the team work it became quite clear to us that the collaboration was not going to be successful. We only managed to get limited contact with two persons from Kenya, and never managed to either skype chat (neither text or voice) or email with the last two members. One thing that could cause their lack of interaction in our e-mail correspondences and skype meetings are the technology readiness. We know that it is a challenge for some of the group members to get internet access. A recent study shows that only 36,3 % have access to internet in Kenya [28], which compared to Denmark's 86 % in 2010 [9], indicates that it is not as easy to get internet access in Kenya as we are used to in Denmark. And we know by fact that a few of them do not have a stable internet connection in their home. Our approach to solve this issue was to have a meeting each Tuesday at 10:00 GMT +1. This way we knew that they should have access to their University, which on normal circumstances has an internet connection they could use. We know they should have time for this meeting, thus it is planned as a course on their schedule. Secondly it might have been caused by collaboration readiness. Collaboration readiness is a potential show stopper for any team work, if a given member is not ready to collaborate. There could be many reasons for this, one being that the student was unsure if he or she was going to follow the course. We have tried to identify these issues towards our fellow group members and we cannot find anything that should indicate that they would not be willing to collaborate. They should be just as interested in delivering a good product. We never received any other indication or statement telling us otherwise.

3.2.3 Tools

Generally we used four different tools to communicate our teamwork and progress; one synchronous and four asynchronous. We used Skype for every team meeting and for general communication between the members of the group. Skype is synchronous communication as you will normally receive an instant reply while using voice chat. We furthermore used four asynchronous communication methods, that is GitHub, Email, Google Drive and Doodle. GitHub was used to share the actual projects current status by communicating via tickets what needs to be done. Email was used for communicating general group announcements, that we want to make sure that all of the members from the group receive. Google Drive to share important documents, such as general to do lists and spread sheet with member interests and skills. In order to mitigate the time zone differences between Nairobi and Copenhagen, that could create problems, we used Doodle to try to sketch up a normal group members week. Doodle is a free online scheduling tool, that automatically takes time zones in to account. One gets presented with a scheme of days and given periods on this day, such as Monday 18th of February, 8:00 to 12:00. Then one selects the hours and days that fits. This was done to sketch up the available hours and days that matched the group best. Four Danish members and one Kenyan member filled out this form. It nonetheless told us, that we would have problems meetings different days than Tuesdays and thus had to rely on the other mentioned asynchronous tools to communicate.

3.2.4 Mitigation

We feel, as a group, that we were well prepared for this collaboration. We only got to use a small subset of our methods to mitigate problems due to the fact that no real intentional collaboration effort was there. It is quite obvious to us that no different tactic or set of methods could have solved these problems. If the collaboration had moved from the initial stage, and reached an actual level of real collaboration, one could probably argue that our approach was relied on wrong methods or the likes. But it failed at the early stages, before we had any real impact. These challenges and our thoughts on what could have been done differently will be discussed and reflected in Section 7.3

3.3 Technology

In the beginning of the project we talked about the many technologies available for web development. Obviously there are many approaches to the same solution. Focussing on working as a team, that all have different background experiences

we decided to individually write down our development qualifications.

From all of the qualifications in the group, we joined together and made a top three technology approached (see Tabel 3.1. Every team member voted 1, 2 and 3 on an development approach. 1 being the most preferred, and three the less preferred.

Table 3.1: Team members experience with different technologies

	Kasper	Thomas	Stefan	Rasmus	Nicolas	Steven	Cecil	Tom	Lucy
PHP	2	1	3	3	1	1	1	-	1
C#.Net	3	3	2	1	3	3	1	-	3
Java	1	2	1	2	2	2	3	-	3

Note that Tom in Tabel 3.1 didn't want to share his experience with the rest of team. Therefore the comparison table was never completed.

One of the challenges in this project was technology readiness. Every team member has a preferred development language, which ultimately led to a long debate on settling with a development language, and what kind of framework the group should use. The problem was that: either way - some team members would have a harder time participating with implementation than others.

Also some members had already lots of experience in development web applications, while other had very little. Not everyone was familiar with MVC based development, which for those unfamiliar became a challenge. The overall problem was simply having different level of knowledge regarding web development and technologies.

Ultimately however the group decided to go with Java as the main programming language, along with Javascript for front-end parts.

Various milestone for the project was planned and the different implementation tasks, was assigned to different team members. Most noticeable the team unintentionally created subgroups: divided in team members working with front-end tasks, and team members working with backend tasks. However, frequent meetings and online conversations led to positive awareness on what everyone was working with.

Furthermore the team utilize an online project management site, from which the team members could keep each other informed of work in progress. It also included a Gantt chart scheduling the entire project development.

Also Github was used to easily share and version code and documentation be-

tween team members. This made it possible for everyone to see changes and follow ongoing development.

Chapter 4

Design

With inspiration from other building management systems (see Section 2) and looking at the building simulator we had to communicate with, we came up with the concept of policies that consist of two types of statements; *if-statements* and *set statements*. If-statements works as if-statements in many popular GPL's¹, with a conditional expression, a then-clause and an else-clause. Set statements work by setting a value on the building simulator (in effect it acts as an actuator).

During the initial research for scripting languages we looked at several established embeddable scripting languages (see Figure 4.1).

Name	Description
Python	Python scripting
JRuby	Ruby scripting
Tcl/Java	Tcl scripting
Rhino	Javascript scripting
BSF	Bean Scripting Framework*

Table 4.1: Established scripting languages found during research

* The BSF supports several languages.

In the team we had some initial discussions if we should use tools that could assist us in designing our own DSL². Examples of these kind of tools could be technologies like Xtext [4], Xbase [3] and EMF [1]. Xtext is a framework for developing programming languages and domain specific languages. Xbase is a partial programming language built in Xtext and used for integrating into other programming languages and domain specific languages. Xbase would be able to

¹General Programming Language (like C++, Java etc.)

²Domain Specific Language

give us an working expression language. EMF is the Eclipse Modeling Framework, a modeling framework and code generation tool. Some of the team members in the danish group had used those technologies during the ITU course Model Driven Development, but no kenyan team members had any experience with these technologies. This would in effect mean that they would need to acquire the skills learned during Model Driven Development, while simultaneous working and contributing to our joint policy engine project. We deemed this unrealistic. Another argument for not using this set of technologies is that Xbase seemed too daunting a project to undertake due to it's current state. At the time of writing, only 25 questions have been tagged with Xbase on [2] - a website used frequently by the danish team members when in need of coding assistance. The lack of Xbase examples and documentation also reveals that it is a very new, and unfinished, product. Since Xbase would be the caretaker of the expression language, we could have opted to forego just that, and still use Xtext and optionally EMF. We considered this but arrived at the conclusion that using Xtext alone would not give us that much, since our language is small. Also the amount of overhead introduced by using Xtext, both in complexity but also in the generated code, would further add to the confusion of our team members with no experience with this technology.

Based on the these arguments we chose to develop the functionality ourselves. We also considered that task to be more *fun*! We believe that developing our own scripting language was best given the resources we had available and the timespan of the project. Considering the embeddable scripting languages (see Figure 4.1) we only needed a small subset of the functionality (mainly the if-then-else construct) and therefore deemed the complexity versus the timewise return of investment unfavorable.

4.1 Architecture

Merriam-Webster defines *policy* as "A definite course or method of action selected from among alternatives and in light of given conditions to guide and determine present and future decisions".

We have defined a policy as a collection of statements operating on sensors and actuators residing in the building simulator. The policy has a start time and a stop time, indicating when it is active. It is also possible to de-activate a policy completely — without deleting it entirely — using a boolean flag.

The system that we have designed consists of four overall parts;

- The management website (front end)




Figure 4.1: The PolicyEngine system architecture.

- The policy engine (back end)
- The database for storing policies (back end)
- The provided building simulator (back end)

Figure 4.1 illustrates the relationship of the system parts. It is evident that communication between both the front end and the simulator primarily is achieved through Json. For a more in-depth explanation, refer to Section 4.1.2.

In terms of MVC³ the *model* and the *controller* is the servlet containing data and business logic. The servlet reacts to several urls (see Figure 4.2), that functions as an interface for an operation needed in the system. The *View* is the rendered html based on JQuery.

A statement can either be a SetStatement or an IfStatement. The SetStatement sets an value in the simulator. The backend implementation allows nested If-statements. If-then-else statements have been around possible before even the early versions of Basic, and is an established way of achieving conditional computer logic. Even though the intended audience of our system is not considered computer experts, the if-then-else concept is easy to understand and catches on fast when interviewing people. An If-statement can contain multiple expressions which are being anded when evaluated. If the user wants to make an If-statement with or'ed expressions, she will have to use a nested if. This raises the complexity and we do not consider this to be an optimal solution. The optimal solution in this case would have been to make a proper expression language that is left recursion safe, for example based on [25]. This would also solve our lacking or's in the condition of the if-statements. Unfortunately we did not realize this until around two thirds into the development. Since our jquery [33] parsing on the frontend naturally is strongly bound to domain classes on the backend (due to JSon serialization using Gson [12]), any big changes at that point would affect

³Model, View and Controller

both the backend and the frontend. We estimated then that it would be too time consuming to change the system. We therefore leave it up to future work, and we elaborate further on this subject in the section 7.

4.1.1 The management website

The management website is a small collection of html pages that relies heavily on JQuery. The building policy administrator opens the main page in a browser, and gets an overview of all the policies and can perform CRUD⁴ operations on them.

During the design phase for the GUI of the management website, we used Photoshop sketches that we mailed and transferred via Skype to each other. One from the kenyan team also assisted in this phase. After discussing how to best present policies to the building policy administrator in the time we had available, we choose to use both indentation and color coding page layout (see Section 5) when presenting the policies. By choosing this type of GUI, as opposed to a more tabular view on the if's, we elected a GUI that is flexible enough to extend to a recursive rendering which clearly matters a greatly when then-clauses can contain other if-statements that again contain then-clauses and so forth. For implementation see Section 5.

4.1.2 The policy engine

The policy engine consists of a servlet and domain classes for the expression language needed to build policies.

A statement can either be a SetStatement or an IfStatement. The SetStatement sets an value in the simulator. The backend implementation allows nested If-statements, making the policies both flexible and simple.

An If-statement can contain multiple expressions that all are being anded when evaluated. If the user wants to make an If-statement with or'ed expressions, she will have to use a nested if. The optimal solution to this would have been to make an expression recursion safe model. We did not have enough time for this, but we will elaborate further on this subject in the discussion (see 7.2).

The statements are organized as a list. Each statement will be executed (set) or evaluated (if) in the order that they are persisted. Each if-statement has two lists that contain the statements belonging to the then-clause and the else-clause.

⁴Create, Retrieve, Update and Delete

The servlet is the core of our policy engine. It has three main objectives;

- The timed, repetitive execution of policies which currently are active
- Serving web service requests to the JQuery front end
- Serving the management website html files

An overview model of the PolicyEngineServlet can be seen in Figure 4.2.

The servlet container used is Tomcat 7.0.37⁵.

This design was based on several iterations of group discussions. Everyone wanted a solution that worked efficient and fast, within the loosely defined scope that we initially had set for the system development. And clearly, our project should fit in the timeframe allowed for the project.

The first couple of design proposals from some team members were based on SQL integration — by the means of creating a java sql connection — directly into the database that the building simulator uses. This direct integration could be argued as being a good, although strongly coupled integration. Also, it did not feel like a real world scenario (which we are sure that ITU intends as being one of the reasons for this course) if we could access the database directly in this way. Traditionally building systems tend to be closed source proprietary systems with few, if any, integration point. Today there exists different kind of protocols (KNX, LON and BACnet to name a few) that allows some integration to existing building automation systems (for example CTS ⁶ systems that are widely used in Denmark). Since our project deals with a building simulator, we do not have to worry about these kind of implementation details — but we are aware that they exist.

⁵Other versions might work as well

⁶Central Tilstandskontrol- og Styling




Figure 4.2: A overview model of the PolicyEngineServlet.

Later design proposals revolved around the idea of having more or less mirrored sensor data in a separate database. Data should be transferred at regular intervals, and the policies should be executed against the copy. This turned out to be a too unsophisticated solution, for several reasons. It would involve a lot of overhead copying data from one database to the other. There is also a problem in knowing exactly what data to copy, since that depends on the statements that is currently active. Based on the above mentioned discoveries we decided on the following design, that takes all team discussions into account and solves all earlier stated challenges.

On servlet container startup the servlet will read various configuration values, and then set up a thread that manages the timed, repetitive execution of active policies. While testing the system policies were executed every 5 seconds, but this is configurable. The servlet will start by querying the database for active policies, which are defined as policies that have a boolean flag (active) set to true, while the current time is within the policie's *from* and *to* time. This will result in a list of policies that should be executed. Then all sensor id values of all if-statements (including nested statements) are fetched from the building simulator, and cached in an in-memory datastructure. Then the policies are executed, and values set accordingly using the rest API.

4.1.3 The database

The database used is mySQL 5.5.29⁷ and used for storing the policies. At the moment we only use one table containing all the policies.

4.1.4 The building simulator

The building simulator [17] provided in the course has not been changed or modified in any way.

⁷Other versions might work as well

Chapter 5

Implementation

This chapter is dedicated to the implementation of the presented system. In this chapter we will explain in detail what has been implemented and how.

5.1 System

The system has been implemented using Java and Servlets for the back-end, and JSP and Javascript for displaying the webpages in the front-end. We have chosen to use MySQL as a database since it is known to all the members. The database is used to store the policy entities (the policy entity is explained in subsection 5.1.2). There is also a logging subsystem that is used to create logs about the system and its execution. This has been implemented using the log4j library and it creates a log file which contains all the logged data.

The back-end deals with requests from the front-end and executes at a specific time interval, the active policies at that time. The front-end displays the policies and enables the creation, modification or deletion of policies. However, because we have chosen to serialize the policies to JSON, in the frontend these policies have to be deserialised and parsed to be displayed properly.

A serialized policy to JSON has the following structure (see Figure 5.1) It contains a statement object which is composed of the objects conditionalExpressions, thenStatements and elseStatements. The conditionalExpressions object describes the IfStatements and has information about the sensor id, the prefix operator, and the operator. The thenStatements and elseStatements have information about the desired value of the sensor and the id of it. For this example, the gain value of the light sensors on the floor 0 are compared to the desired value, 1 with the help of the "EQUALS" operator. If the gain value of a sensor is equal to 1, then the light is turned of by setting its gain to 0.

```

{
  "statements": [
    {
      "type": "dk.itu.policyengine.domain.IfStatement",
      "data": {
        "conditionalExpressions": [
          {
            "prefixOperator": "AND",
            "aValue": {
              "type": "dk.itu.policyengine.domain.FloatValue",
              "data": {
                "floatValue": 1.0
              }
            },
            "operator": "EQUALS",
            "sensorId": "wildcard-floor-0-light-gain"
          }
        ],
        "thenStatements": [
          {
            "type": "dk.itu.policyengine.domain.SetStatement",
            "data": {
              "aValue": {
                "floatValue": 0.0
              },
              "sensorID": "wildcard-floor-0-light-gain"
            }
          }
        ],
        "elseStatements": [
        ]
      }
    }
  ]
}

```

Figure 5.1: A serialized policy




Figure 5.2: The core classes in the *expression language*.

5.1.1 Interfaces

There are two interfaces that describe the methods of Statements and Value classes. Each statement should implement an execute method, in which specific execution is handled.

5.1.2 Domain

A statement can either be a SetStatement or an IfStatement. The SetStatement sets an value in the building server (in effect it is an actuator). It is possible to have nested IfStatements, making the policies both flexible and simple. An IfStatement can contain multiple expressions that all are being anded when evaluated. If the user wants to make an IfStatement with or between the expressions, will have to use a nested IF. The optimal solution to this would be to make a safe left-recursive model. We did not have enough time for this, but we will elaborate further on this subject in the Discussion section.

An Expression has three variables, the sensorId, the operator and a value of type Value. When the policy engine is running, each expression is being evaluated. The current value of the sensor is fetched and compared to the desired value using the operator. An expression can contain the following operators; `<`, `>`, `<=`, `>=`, `!=`, `=`. If the evaluation of the expression is true, then the statement is executed.

The FloatValue class implements the Value interface and it represents a value of type float. This is used to set the values of the sensors, as those can be of float type. However, at this moment we only support to set integer values, 1 for On and 0 for Off.

A Policy object consists of statements and the policy can be run by executing those statements. A PolicyEntity object holds a policy and different properties of that policy such as name, description, interval, id and active. The interval property defines from which time to which time should the policy be enforced to

run. This is implemented as a separate class, `Interval`, because of the serialization logics. The active property represents the status of the policy as in if the user wishes for this policy to be executed (active) or not (not active).

To be able to handle many sensors and actuators, we have implemented Wildcards which define the sensors and actuators handled, and their type property (gain,setpoint). An example of a wildcard is "wildcard-floor-0-light-gain". This expresses that we want to retrieve information about the gain of the light sensors from the floor 0. The `ThenStatements` and `ElseStatements` can be executed using the same wildcard, a different one or just some different sensors.

5.2 Persistence

We have decided to store our policies in a database as we think it is better to have them on a different server than the one running our policy engine. This may be a drawback if there if the internet connection fails to either servers. However, volatility is always a problem with distributed systems and this issue could be handled by having a backup database that resides on the policy engine's server and is synchronized with the main database server.

The database chosen was MySQL as it is well known and used by the group. Because we do not have a complex persistence system, we decided to use the simple JDBC for connecting and querying the database, and not use other specialized frameworks such as Hibernate¹.

Methods for communicating with the database are defined in the `DataAccess-Layer` class. These are the methods for creating, reading, updating and deleting policies. In addition, it is of great interest to have a method that returns only the policies that are running at the current time and are active.

Different policies may operate on the same sensors, so we decided to use a cache to store each sensor's value. This reduces the number of queries to the building server so it does not overload it. This cache is implemented using a hashtable and stores the value of the sensor. After executing all the active policies at that moment, this cache is cleared.

5.3 Building querying

To communicate with the building's server, we created the `Connection` class which provides methods for connecting to it and parses the JSON object received. There are several methods implemented that parse the response of the server and offer information such as - a list with the sensors in a room by providing the room id,

¹<http://hibernate.org/>

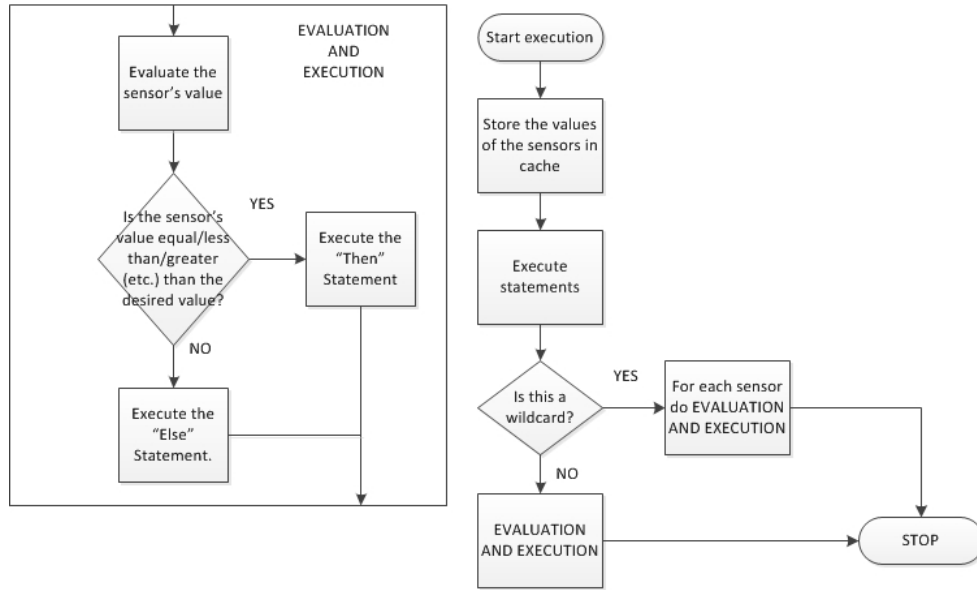


Figure 5.3: The execution flow of a policy.

a list with all the sensors of a certain type (e.g ac, light, heater etc.) The two most important methods are for retrieving the sensor's value and setting it to a specific one.

To retrieve a sensor's value we use the `GetSensorValue` method which takes as an argument the sensor id and returns a string with the value. By default the values are sorted by timestamp, and the most recent value is returned.

To set a value we use the method `SetSensorValue` which takes as parameters the sensor's id and the value we want to set.

5.4 Executing policies

The system executes the active policies at a specific time interval. This execution has the following flow 5.3: First, a list with the active policies at that time is retrieved. Then the statements that compose the policy are retrieved along with their expressions. In these expressions resides the sensor id on which operations will occur next. For each sensor it is retrieved its last value from the building server and stored in the cache. After storing all the values of the sensors, the statements are executed.

The first statements that are executed are the `IfStatements`. Here a check if the policy is a wildcard is done. If it is a wildcard, then the statements are executed for each sensor by evaluating its value against the desired value. If the

evaluation returns false, then the else statements are executed. If the policy is not a wildcard, an evaluation of the sensor's value it's done and depending on its result, the then statements or else statements are executed. In the evaluation part the actual value of the sensor is compared to the desired value using the operator defined in the expression. This evaluation returns true or false and so execution of the thenStatements is done if the evaluation is true otherwise the elseStatements are executed. In the execution of the then- and else-statements, the new sensor value is set in the building's server.

5.5 Front-end

5.5.1 Initial Ideas

We quickly realised that our ideas could easily expand the project, with numerous functionality and features, that would bring the project to a far too complex level. We mapped all our ideas and prioritized them, and made a selection of ideas, which we planned on implementing.




Figure 5.4: Initial drawing.

The selected ideas that we wanted to realise (sorted by importance) included:

- Users should be able to add, delete and modify policies.
- Users should be able to use complex operators in policies, e.g. higher than, lower than, equal to etc.
- Users should be able to save policies and easily toggle an ON or OFF state.
- Users should be able to combine multiple sensors in a single policy.

- Users should be able to use wildcards in a policy, e.g. effecting for instance an entire floor without the need to specify the rooms that belong to the chosen floor.
- Users should be able to create nested rules in policies.

Policies
Add, delete or modify policies

<input type="checkbox"/>	FOR	Floor 3	Step 1								
<input type="checkbox"/>	FOR	Floor 3	Room 2	Step 2							
<input type="checkbox"/>	FOR	Floor 3	Room 2	IF	Temp	IS > 20	Step 3A				
<input type="checkbox"/>	FOR	Floor 3	Room 2	IF	Temp	IS > 20	Step 3B				
				AND	Whether	IS = Sunny					
<input type="checkbox"/>	FOR	Floor 3	Room 2	IF	Temp	IS > 20	THEN SET	Aircondition	=	ON	Step 4
<input type="checkbox"/>	FOR	Floor 3	Room 2	IF	Temp	IS > 20	THEN SET	Aircondition	=	ON	Add
				AND	Whether	IS = Sunny					

DONE

Figure 5.5: Initial drawing - step wise build up.

All these points were realised in the backend of our implementation, however because of time limits we regret that the possibility to create nested statements did not make it to the user interface. Still we have an interface fully ready to support nested statements, but this is something we have added to future work instead. We have started work on a recursive function to be able to cleverly support endlessly (in theory, computing resources are not endless) deep nested statements. These nested statements would be listed under THEN and would contain a new set of IF, THEN, ELSE subsections. These sections would be easy readable because of indentation and encapsulation.

5.5.2 User Interface

The main goal for the user interface has been to keep it simple and user-friendly without losing advanced functionality. Initially we started making rough drawings to better understand what we were dealing with and how to best possibly present it. Doing so we tried to sketch the process of handling a policy from the creation to the enabling of it. We came up with various approaches one of the more interesting ones being the step wise (see figures 5.4, 5.5). However this approach proved to not be as flexible and able to handle very detailed and complex policies. Our aim has been to make a great overview for the user. Therefore we chose another concept that is splitting each statement up into an open / closeable area to better encapsulate the logic. We found that you very often read one statement at a time, therefore you can have the statements open you are interested in, and won't be confused by statements not applicable to an area of interest. Within each statement we also decided to split up IF, THEN and ELSE into subsections because you in theory can add as many entries to each subsection as needed. Nonetheless the introduction of "wildcard" values for instance allows to manipulate an entire floor's heaters, taking up much less space and work for the user to define then having to set up entries for each heater on a floor. See an example from the final view of a policy build at figure 5.9.

5.5.2.1 Intuitive Design

What you want to achieve when designing an interface is to make it as intuitive as possible. This can be achieved by looking at the user's "knowledge space" [5].

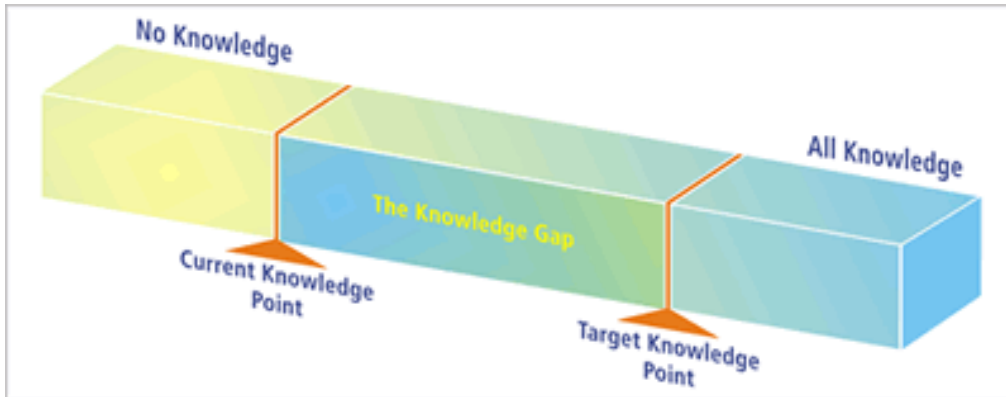


Figure 5.6: The space between the Current Knowledge and Target Knowledge points is called The Knowledge Gap.

The distance from the left (As seen on figure 5.6) indicates how much a given user

is familiar with an interface, this is called the "current knowledge point". The next point of interest is called "target knowledge point". This section details how much knowledge the user needs to know to achieve their goals. Each time a user attempts to perform a specific task, their "current knowledge point" and "target knowledge point" are very important for us to identify. Each user's "current knowledge point" and "target knowledge point" will be different when they gain more experience. However, it turns out that, by making usability tests we can often identify groups of users who have almost the same "current knowledge point". This is also why we have performed think-aloud tests, which we will get back to later in this report.

"The Knowledge Gap", which is the distance between the "current knowledge point" and the "target knowledge point", is also called "The Gap". "The Gap" is what you have full focus on designing for. There is no need to design to the left of the "current knowledge point" and you do not need to design to the right of the "target knowledge point" since the user does not need the knowledge to perform the task. One should therefore focus only on designing the interface for the space in between the two points.

Users can accomplish their goals when the "current knowledge point" is equal to the "target knowledge point". There are two ways to achieve this. You can train the user, thereby increasing their current knowledge, until they know everything they need to do the job. The other way is to design an interface that can reduce the "current knowledge point" by making the interface more accessible until the "target knowledge point" requires only the information the user already has.

5.5.2.2 Flow

We have strived to create an interface that provides a good flow to the user. The user should not notice the interface but have full focus on the task to be solved - the interface should be transparent. Any distraction breaks the flow. [6]

"No matter how beautiful, no matter how cool your interface, it would be better if there were less of it." - Alan Cooper

There are three prerequisites for achieving flow:

- The user must have a clearly defined goal.
- There must be a good balance between the user's skills and challenges the user must deal with.
- There must be clear and immediate feedback.

Especially the last point of the 3 above can be difficult to achieve, when you are in an online environment with response times. However, we believe that we have solved this by allowing the user to model their policy without any server interaction until they want to save.

5.5.2.3 Response Times

In order for the user to get the best possible experience on a website, it is important to optimize response times. Below is the overview of user's reactions at different response times from Alan Cooper - About Face 3:

Up to 0.1 seconds	Users perceive the systems response to be instantaneous
0.1 – 1.0 seconds	Users notice a delay, but their thought processes stay uninterrupted
1.0 – 10 seconds	Users regard the system as slow, and their mind is likely to wander
After 10 seconds	The system will loose the users' attention. They start doing other things or switch to another application

Based on this it is important to optimize the web site, both for the code generated on the server, but also to optimize the front-end since around 80% of the load time [31] of a website is spend on parsing the generated HTML document, execute javascript and download any external resources. There are a number of things you can do, which we will not go in detail with but just touch on. A few points we are aware of is where and in what order javascript and css is located and loaded, javascript can be loaded last as javascript usually is loaded after a page is fully parsed. Furthermore, you can set preferences for cache and reduce the number of external files to load in an application by combining various javascript files into one and also combine CSS files.

5.5.2.4 Color Coding

Color can be a powerful tool to improve the usefulness of an information display in a wide variety of areas if color is used properly. Conversely, the inappropriate use of color can seriously reduce the functionality of a display system [30].

Colors send signals to the user sub-consciously, so therefore you can make use of this with things like buttons and colored areas. We have used this idea to color the different boxes under a statement, to improve the readability and make it

more intuitive. We used a red color for the IF boxes and a green color for the THEN and ELSE boxes. The reason for this is that the IF is a "get" type and the other two is a "set" type.

5.5.2.5 Responsive Design

Responsive web design is a web design approach aimed at crafting sites to provide an optimal viewing experience—easy reading and navigation with a minimum of resizing, panning, and scrolling—across a wide range of devices (from desktop computer monitors to mobile phones). While we haven't directly optimized for mobile we have had the responsive design idea in mind when setting up the design. What we have assured is that the user interface is adjusted to the screen size of the user. Meaning if you for instance use a full HD resolution (1920x1080) the interface will utilize this space (see examples on figures 5.7, 5.8, 5.9).

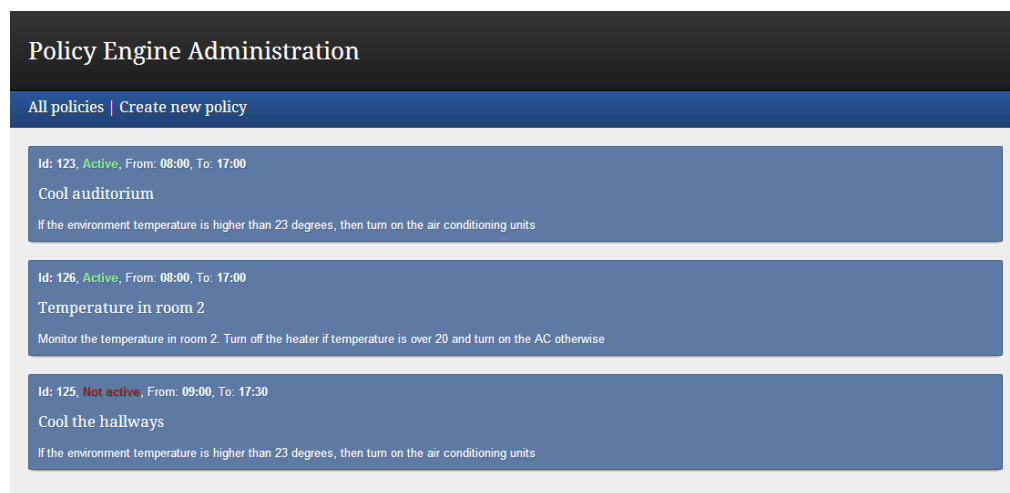


Figure 5.7: Illustrating the policy list-view.

5.5.3 Policy Editor

5.5.3.1 Managing Policies

In order to make it simple and fast to manage a policy we have split the interface for a policy up into sections and sub sections. We have two sections: Policy Information and Policy Statements. In Policy Information, all the content: name, description and time interval are defined. Also if the policy should be active, and run on the Policy Engine, can be defined in this section. The Policy Statements section is where all the conditions and actions are defined. We have added a sub

section for each statement to improve the readability and overview of what is executed and what is paired together (see figure 5.8).

Policy Information (id: 123, Not active)

Is active: ☒ Yes ☐ No

Name: From: To:

Description:

Policy Statements (+ New)

Statement 0	X
Statement 1	X
Statement 2	X
Statement 3	X
Statement 4	X
Statement 5	X
Statement 6	X
Statement 7	X

Update

Figure 5.8: Editing a policy.

Within a statement there are three sub sections: IF, THEN, ELSE (see figure 5.9). For each headline in this setup there is a "+ New" button that can be clicked. This will add a new entry to this particular area. While this is done we alter the policy object in the background so once the user are satisfied with the setup of statements, it can be saved with a click on the update / save button.

To help the user find the correct sensor / wildcard option, we have implemented an autocomplete functionality. This lets the user start typing a value and then the autocomplete will come up with all the matching elements available.

To make it take less space and work to define a statement that should apply to an entire floor or room, we have introduced wildcard options. A use case can for instance be: A user wants to work with the blinds on the entire first floor: floor1-blinds is then the option to pick.

5.5.4 Use Case View

The solution's final use case diagram is shown in Figure 5.10.




Figure 5.9: Configuring statements.

Figure 5.10: Use Case Diagram (that is missing!)

5.5.5 Usability

Seeing policies as rules for rooms in a building, with multiple sensors in each and every room, that can affect multiple actuators - not only in one room, but all of the rooms in a building: It leaves the user with a lot of selections. This will, if not represented properly, complicate the process of adding new policies.

Following the principles of Steve Krug's "Don't Make Me Think: Common Sense Approach to the Web"[22], the functionality of a website should always let users accomplish their intended tasks as easy and directly as possible. Throughout the development of the front end site we have thrived to do so. And we have had two usability tests (more on this in the Evaluation part) to uncover potential problems.

5.5.5.1 Browser Compatibility

We have optimized for modern web browsers. The specific browsers we have tested our solution in are: Internet Explorer 9 and 10, Firefox, Chrome, Safari.

5.5.5.2 Layers of Abstraction

With the policy engine there are different layers of abstraction: Visual UI, Textual UI, Engine.

We have focused on the Visual UI enabling a graphical representation of the CRUD process. The other approach is having policies handled in a textual manner, and thereby forcing users to write the statements in a policy, using a domain specific language. For a developer textual editing and console commands might be preferred, and may somewhat be quicker, while they know the inputs by heart. But the intended users of our policy engine, are most likely not developers or fans of typing in commands.

In a future version of the policy engine, both methods could be implemented to function in parallel, giving a user both choices, and also the capability of copy and paste policies quickly.

Chapter 6

Evaluation

In our project we have two areas that we have been evaluating. The first being whether our policy engine is actually working and enforcing the defined policies. The other being the front-end interface, testing to uncover potential usability issues that needed solving. This is important as we need to see how well our system performs and to see what has been implemented from the requirements.

6.1 Policy Engine System Evaluation

6.1.1 Log Testing

We have logged system data which was important as to determine if the system is working correctly. Doing so we have been able to verify the results and outcome while testing and to discover bugs in our implementation. For quality control we have also been using these log files. Verifying our solution by cross-referencing every action during tests, and match it with the actual data stored in the database. One excerpt from the log file can be seen in figure 6.1

```
|2013-05-11 23:56:28,122 [Thread-1] INFO dk.itu.policyengine.domain.Expression - value is 0.0 -  
|2013-05-11 23:56:28,122 [Thread-1] DEBUG dk.itu.policyengine.domain.Expression - Is 0.0 equal to 0.0 ?  
|2013-05-11 23:56:28,122 [Thread-1] DEBUG dk.itu.policyengine.domain.Expression - Yes!  
|2013-05-11 23:56:28,122 [Thread-1] DEBUG dk.itu.policyengine.integration.Connection - Trying to connect to  
|http://gsd.itu.dk/api/user/building/entry/set/1/room-14-heater-14-gain/1/?format=json  
|2013-05-11 23:56:28,242 [Thread-1] INFO dk.itu.policyengine.integration.Connection - {"returnvalue": true}
```

Figure 6.1: An excerpt from the log file.

6.1.2 JUnit Testing

To further test the software we made, we performed a range of JUnit Tests. After the creation of the domain classes we have created JUnit test to verify if everything works as expected. One example of these test can be seen in ???. In this test we set the temperature of a room to 11 degrees and we want to turn on the heater in the room if the temperature is equal to 10 degrees, otherwise we want to roll down the blinds. The test will succeed when checking if the blinds are rolled down, as the temperature is set to 11 degrees. We have employed tests for checking all the possibilities of operations against the operators. The result was that our IF-THEN-ELSE concept implementation performed as expected.

6.2 Usability Test

To evaluate on the usability of our policy engine we decided to make a test with candidates outside of our development group.

In general, when it comes to "best practices of usability tests", the Think Aloud Protocol is considered one of the most valuable [26].

Therefore we decided to use it for our evaluation. Also the Think Aloud Protocol, is inexpensive in cost and time, and it is easy to set up and gives very valuable results from real-life scenarios.

6.2.1 Think Aloud Protocol

In a thinking aloud test, you ask test participants to use the system while continuously thinking out loud — that is. Simply verbalizing their thoughts as they move through the user interface, and take actions.

To run a basic thinking aloud usability study, 3 things are required:

- Recruit representative users.
- Inform them of representative tasks to perform.
- Avoid interference and let the users speak their actions.

We invited five people to a think aloud test. Research shows that having just five people will potentially uncover 80% of all usability problems [27] as seen on 6.2.




Figure 6.2: Usability test graph.

The candidates is male and female students at ITU, between the age of 25 and 30. Some of them study software development and have a lot of knowledge in programming, others study digital design and knows more about human computer interaction and usability. However they all have web development in common.

6.2.1.1 Tasks

We arranged 7 tasks for the candidates to perform. Note however that the tests where held individually from each candidate, but they were all giving the same set of tasks. During the tests, candidates where only given one tasks at a time to focus on - using small tasks cards.

The candidates where assigned the following tasks. The questions is designed so that the candidates makes use of all the features in policy engine:

Before the test started we explained the purpose of the Policy Engine, and verbally gave them an example of a user case scenario. We briefly introduced the participant with a map of the building and the rooms involved. We also gave an example list of sensor and actuator -names to get them familiar with the elements involved including basic knowledge of IF, AND, THEN statements and wild card operators.

Note that the user interface is designed with "live auto complete functions" so

that the user can easily set a desired sensor/actuator without knowing the exact name. (SEE FRONTEND DESIGN SECTION)

Create a policy that turns the AC on in Room 1, 1. floor and name it "Cooling".

Modify the "Cooling" policy you just created to affect both Room 1, 1. floor and Room 2, 2. floor.

Create a policy that turns on Blinds in all rooms for every floor in the building and name the policy "Sun".

Find and show the active policies just created.

Disable the policy named "Cooling" so that is no longer in function.

Delete the policy named "Sun".

To save energy the university wants to have the heating turned OFF automatically at 17:00. However this Wednesday around 19:00 - 22:00 an exclusive presentation is held in room number 5 on 1. floor. You are asked to maintain a temperature at 21 degrees in that room throughout the presentation.

It is summertime and the overall temperature inside the building is rising. You are asked to keep temperature at maximum 22 in all of the rooms in the building.

All afternoon between 12:00 and 16:00 the sun is at its peak. Therefore you are asked to set blinds ON in all the rooms on 1. floor and 2. floor - but only if the lights are ON.

The university wants to save energy. You are asked to make sure lights are automatic turned OFF at 17:00 in all the rooms, except from those on 0. Floor.

6.2.1.2 Results and Comments

After every participant had gone through the think aloud test. We went through all their remarks. Some were the same, and those that were - we combined into one. For all the remarks, we made some comments as listed below - including the actions we took to further improve the software.

If there is a lot of policies - which I would expect there will be?
Then I think the active policy list will be too long. I think it would make it difficult to find a specific (policy) if say you have a 100 policies.

The user is referring to the “Active Policy Site”, from which every active policy is listed underneath each other. The user finds that the site may force the user to do a lot of scrolling in order to find a specific policy.

Due to this we implemented a expand function, which instead of listing all the content, in all of the listed policy. It on shows name and time and a button in the right hand corner, which if clicked expands the policy and reveals all of the content. This way the list will be much more compact and easier to go through.

For future implementation we would also like to make a search function, from which the user is able to filter the policy by the search inputs. In example: name, or rooms effected by the policy etc.

I think it is difficult to see which one (rules) belong to what (expression).

The user is referring to the rules within a given policy. In every policy the user have the option to use expressions such as IF, THEN, ELSE. The user finds that it is difficult to see which rules belong to what expression (IF, THEN and ELSE).

Due to this we implemented a color technique, which colorize the rules that belongs under the same expression. As an example all the rules belonging to an IF -expressions is colored with a red background on the site, while all the THEN -expression is colored green.

In the second think aloud test, all participant expressed that they loved the coloring. That they diffidently feel that the overview is much better and think it is easier to understand.

6.2.1.3 Subset

After the first think aloud test, we made changes to front-end design and we repeated the think aloud test, using the same set of tasks, to verify the improvements.

After the second round of testing we only had some minor changes to do. But overall the participant where comfortable with usability and had no further remarks. We believe that there is always room for improvement. And so we have made some comments on this in section 7.2.1.

Chapter 7

Discussion

In this section we are going to discuss our project in three different parts:

- The collaboration with the team members from Kenya
- The prototype system
- The overall project

We have chosen to analyze each of these areas with an approach of highlighting what could have been done differently and perhaps in a better and more successful way. We will reflect on our choices made during the project and finally also on our learning outcomes.

7.1 Collaboration

The collaboration proved to be a rather challenging task. Only one person from Strathmore showed interest in the project, even though this should have been four persons. We believe that this is not caused by our approach towards the students from Kenya but either; 1) a lack of interest and willingness to work on the project or 2) bad luck with the student selection - they might be swamped in work or having too many demanding courses.

We believe we did all we could in regards to motivating our foreign colleagues. In the beginning we sent them an e-mail to introduce ourselves, based on the TA's experience regarding past collaborations with students from Kenya. We asked them in a kind way to update our "contact information" document, so we could reach each other on Skype/chat/email etc. The last member from the Kenyan

group updated the contact information after 16 days - and thereafter was not heard from again. We created a mailing list (gsdall@netstarter.dk) with all our email addresses to ease the communication between all the members. All the emails were written in english, even when collaborating in the danish group. We wrote out status messages, at least once every week. During the entire project, we only heard from two of the kenyans. One of them counted for almost all communication. We also tried to schedule Skype meetings, on our weekly full working day (tuesdays). However, only one kenyan was there, and he was only online around half the tuesdays. And often it was very late (16-20:00) until he came online, while the danish team members had been at it since 10:00 AM (12:00 PM in Kenya). Perhaps this would have worked better if we would have switched time for the meetings, but unfortunately this was not possible due to time overlapping with other courses.

In order to improve our collaboration, it might have been possible to create subgroups, ie. one from Kenya and one from Denmark - and then assigned those groups a task. These subgroups would have been useful for the implementation part, where different tasks would have been assigned to subgroups, making it easier to communicate and . Because the fact that only once in a while there were two members of the kenyan group online, we decided to make two subgroups in our team, one working for the front-end and one on the back-end, while keeping informed the kenyan group of the progress. When we could begin to see the deadline approaching, we only involved that one active guy from Kenya in some GUI developments. The others would have been invited to participate also, clearly, but they never came online or answered any of our 30+ status emails. At one point, we tried to discover what are the issues in the Kenyan group for not joining in the project and we were told that either they do not have time or it is hard to get online. Also, at this moment one of the members said that it would not participate at all due to personal reasons.

7.2 Prototype system

Our prototype system has implemented the requirements to communicate with the building's server and execute policies. The user interface allows the possibility of creating, deleting, updating a policy and visualizing all policies. Details about one policy can be seen, and updates can be made. We have supplied several JUnit test classes to create policies we find relevant. When policies are being made and persistent, they *are* executed against the live building simulator. Therefore we believe that the system, as a prototype system, is working sufficiently well. Though we have tried to make the policy engine as flexible as possible, some issues remain. For example, if we are having a policy that rolls down the blinds during the night, they should roll up when the policy is no more enforced to run

(decided by the `toTime` property). This is not possible at the moment just by having one policy. We can of course create a new policy that rolls up the blinds at a certain time in the morning, but this is not the best way to handle this issue.

7.2.1 Future Work & Improvements

As explained earlier, our expression language should be improved to allow multiple expressions with the choice of OR as opposed to the current AND-only functionality. With the GPL's on the market today, people expect to be able to make complex conditions - even though the users of this system are not IT experts.

There should also be made some enhancements to the `SetStatement` functionality. We have discussed several functionalities that would be nice to have, but the two most important ones are *Time* and *State*. Explained shortly, *Time* can be used in two ways. It should be able to conditionally query if the current time is within some specified range. This could be used for making policies behave differently on certain times (for example, from 14:00 to 19:00 on Wednesdays). Time can be added relatively simple, as a basic-edition wrapper of the `GregorianCalendar` class. *States* are thought to be boolean variables accessible globally and locally - dependent on which scope it was declared. This way, for example, a fire detection policy could set the *buildingIsOnFire* state, and each other policy would be able to change it's behavior if needed based on access to that state. States can also be programmed easily, by having a map or a hashtable with key and boolean value pairs. The local states can be implemented by using the same structure for the global states, but with the policy id as a prefix in the name.

The reason we did not develop the above mentioned functionalities was the time pressure.

7.3 Project

In this section we want to discuss the overall project and not only focus on the end product. Instead we want to put the collaborative experience in the spotlight and focus on what could be improved regarding the collaboration, such as tools and methods.

It has become clear to us that we could not have a lot of things differently in regards to the collaboration. This has become clear because we believe that the lack of interest in the project from the Kenyan group does not originate from our approaches. Two members have directly responded that they do not

wish to participate in the project, one is not actively participating and the last group member is highly unreliable. These motivational challenges are core values that one needs to have. No strategies or collaborative approaches could have transformed the situation to something positive. These challenges are not a part of a normal work setting, thus one probably would be reprimanded by not attending meetings, replying to emails and show lack of motivation.

If the initial motivation was there, a set of different methods could have been used to kick off a great collaborative team. Some of them are discussed in 3. One approach that could have been useful was to introduce each others countries and talk about topics not directly related to the core project. We actually tried opening up for this kind of communication early in the project, by making a "Swahili - Danish" dictionary document. The active kenyan team member helped filling out different words, and the idea was then that we should try to speak a little in each others language when Skyping. Unfortunately, due to the general lack online kenyan team members, we never really used it. The few times we actually communicated with them via voice, it would have been too 'forced' and strained to use that approach. We would have liked to communicate more and more often with our kenyan team members. This way one would learn the different team members to know, and know more about them than their professional skills. This could have been done by just simply introducing one self, family life, interests, hobbies and the likes.

We suffered a lot from this during our especially first iteration. Our hopes and expectations for the collaboration was high, and thus we wanted our Kenyan team members to give us feedback and inputs on our early choices of programming platform, setup and the likes. We probably waited too long time for their opinions, which caused a delay in our project schedule. We could have solved this by creating a group contract that one had to obey, and if not, one was not part of the group any more. This was discussed during the danish team members, but it would have been a ultimatum - put forth mainly due to frustration and not to improve the collaboration. We could have used more strict deadlines and structured the work even more. This way we would have an early indication of their motivation, or lack hereof, and would be able to push the project forward without their involvement.

We decided to assign different switchable roles to each person in the group. The roles switched during the project, but our focus was to achieve the highest performance throughout the process. The roles were assigned through an analysis of each persons set of skills, motivation and interest. Some persons focused more on the backend, while others worked on the frontend, and others on the report, as their primary tasks. This has showed to be a useful tactic because of two different reasons: First of all we know that the assigned person is motivated to work with the assigned area and secondly has the skills, or wants to learn how,

to solve the problem at hand. This should not be interpreted as a person was 'stuck' doing only one thing. Team members could easily shift to other areas, and that did also happen.

Another useful tactic was our weekly meetings, with a simplistic agenda: What is the current status and how can we push the project forward? All of our communication during these meetings were very specific and all related to these challenges and how to solve them.

It should have been possible for us to directly contact the assigned Kenyan teacher, and asked him to contact the students and have them explain to him what hinders them to participate. If ITU are going to do a collaborative project with Strathmore University again, we would advice that ITU implements much tighter teacher to teacher communication - resembling employer to employer conversations held in private companies when venturing into a partnership. We believe it should have a consequence that a 5 man team effectively is only a unreliable 1 man team. We expect that the Global Software Development should somewhat resemble real world setting. We find it unrealistic implementing such an idea without any leverage mechanism.

We also believe that it could have made a change if the Universities coordinators agreed upon a specific kick off date, where the team members would meet online and coordinate the work from here. This would ensure that we at least got to meet all of our team members.

Chapter 8

Conclusion

In this report we showed how one can implement a policy engine and that our approach to this is valid. We furthermore explains what steps were taken in order to collaborate between two different countries and ten different people.

NOTE:Any cool ideas for a conclusion...?

Bibliography

- [1] Eclipse modeling framework. <http://www.eclipse.org/modeling/emf/>. Online; accessed 20-April-2013.
- [2] Stack overflow. <http://stackoverflow.com/>. Online; accessed 20-April-2013.
- [3] Xbase wiki entry at eclipse.org. <http://wiki.eclipse.org/Xbase>. Online; accessed 20-April-2013.
- [4] Xtext home. <http://www.eclipse.org/Xtext/>. Online; accessed 20-April-2013.
- [5] What Makes a Design Seem 'Intuitive'? Internet. http://www.uie.com/articles/design_intuitive/, January 2005. Online; accessed 10-May-2013.
- [6] Alan Cooper. About Face 3: The Essentials of Interaction Design. 2007.
- [7] D. Dietrich, D. Bruckner, G. Zucker, and P. Palensky. Communication and computation in buildings: A short introduction and overview. Industrial Electronics, IEEE Transactions on, 57(11):3577–3584, 2010.
- [8] Catherine Durnell Cramton and Pamela J Hinds. Subgroup dynamics in internationally distributed teams: Ethnocentrism or cross-national learning? Research in organizational behavior, 26:231–263, 2004.
- [9] Folketingets EU-Oplysning. Internet. <http://www.euo.dk/fakta/tal/internet/>, August 2011. Online; accessed 7-May-2013.
- [10] Facebook. Green of facebook. <https://www.facebook.com/green>, April 2013. Online; accessed 2-May-2013.
- [11] Cristina B Gibson and Mary E Zellmer-Bruhn. Metaphors and meaning: An intercultural analysis of the concept of teamwork. Administrative Science Quarterly, 46(2):274–303, 2001.

- [12] Google. Gson. <https://code.google.com/p/google-gson/>, April 2008. Online; accessed 2-May-2013.
- [13] Google. Google green. <http://www.google.com/green/>, April 2013. Online; accessed 2-May-2013.
- [14] Jonathan Grudin. Groupware and social dynamics: eight challenges for developers. Communications of the ACM, 37(1):92–105, 1994.
- [15] W.B. Gudykunst. Cultural variability in communication. Communication Research, 24(4):327–348, 1997.
- [16] Dae-Man Han and Jae-Hyun Lim. Smart home energy management system using ieee 802.15.4 and zigbee. Consumer Electronics, IEEE Transactions on, 56(3):1403–1410, 2010.
- [17] ITU. Gsd2013 building simulator. <https://gsd.wikit.itu.dk/GSD2013+Building+Simulator>, April 2013. Online; accessed 11-May-2013.
- [18] Rod Janssen. Towards energy efficient buildings in europe. n/a, 2004.
- [19] Sirkka L Jarvenpaa and Dorothy E Leidner. Communication and trust in global virtual teams. Journal of Computer-Mediated Communication, 3(4):0–0, 1998.
- [20] Peter Johnson-Lenz and Trudy Johnson-Lenz. Post-mechanistic groupware primitives: rhythms, boundaries and containers. International Journal of Man-Machine Studies, 34(3):395–417, 1991.
- [21] Andrew Krioukov, Gabe Fierro, Nikita Kitaev, and David Culler. Building application stack (bas). In Proceedings of the Fourth ACM Workshop on Embedded Sensing Systems for Energy-Efficiency in Buildings, BuildSys ’12, pages 72–79, 2012.
- [22] Steve Krug. Don’t Make Me Think: A Common Sense Approach to the Web (2nd Edition). New Riders Publishing, Thousand Oaks, CA, USA, 2005.
- [23] Piritta Leinonen, Sanna Järvelä, and Päivi Häkkinen. Conceptualizing the awareness of collaboration: A qualitative study of a global virtual team. Computer Supported Cooperative Work (CSCW), 14(4):301–322, 2005.
- [24] Jian Li, Jae Yoon Chung, Jin Xiao, J.W. Hong, and R. Boutaba. On the design and implementation of a home energy management system. In Wireless and Pervasive Computing (ISWPC), 2011 6th International Symposium on, pages 1–6, 2011.

- [25] Robert C. Moore. Removing left recursion from context-free grammars. In Proceedings of the 1st North American chapter of the Association for Computational Linguistics conference, NAACL 2000, pages 249–255, Stroudsburg, PA, USA, 2000. Association for Computational Linguistics.
- [26] J Nielsen. Iterative user interface design. IEEE Computer, 26(11):32–41, 1993.
- [27] Jakob Nielsen. Why you only need to test with 5 user. <http://www.nngroup.com/articles/why-you-only-need-to-test-with-5-users/>, March 2000. Online; accessed 7-May-2013.
- [28] Evelyn Njoroge. 14m kenyans have access to the internet. <http://www.capitalfm.co.ke/business/2012/01/14m-kenyans-have-access-to-the-internet/>, January 2012. Online; accessed 7-May-2013.
- [29] Gary M Olson and Judith S Olson. Distance matters. Human-Computer Interaction, 15(2):139–178, 2000.
- [30] Part 1: What the 80/20 Rule Tells Us about Reducing HTTP Requests Performance Research. Internet. <http://www.yuiblog.com/blog/2006/11/28/performance-research-part-1/>, November 2006. Online; accessed 10-May-2013.
- [31] Microsoft Publishing. Microsoft environment. <http://www.microsoft.com/environment/>, April 2013. Online; accessed 2-May-2013.
- [32] John Resig. jquery. <http://jquery.com/>, January 2006. Online; accessed 2-May-2013.
- [33] Kjeld Schmidt and Liam Bannon. Taking cscw seriously. Computer Supported Cooperative Work (CSCW), 1(1-2):7–40, 1992.
- [34] Philip Haves & Thierry Stephane Nouidui Simulation Research Group: Michael Wetter. Building controls virtual test bed. <http://simulationresearch.lbl.gov/projects/bcvtb/>, April 2013. Online; accessed 2-May-2013.
- [35] J. M. Taylor and G. M. Murch. The effective use of color in visual displays: Text and graphics applications. Color Research and Applications, 11(3):3–10, 1986.
- [36] A. Tiberkak and A. Belkhir. An architecture for policy-based home automation system (pbhas). In Green Technologies Conference, 2010 IEEE, pages 1–5, 2010.

- [37] Marcie J Tyre and Wanda J Orlikowski. Windows of opportunity: Temporal patterns of technological adaptation in organizations. Organization Science, 5(1):98–118, 1994.
- [38] Mary Beth Watson-Manheim, Katherine M Chudoba, and Kevin Crowston. Distance matters, except when it doesn't: Discontinuities in virtual work. n/a, 2007.
- [39] Yao-Jung Wen, Dennis DiBartolomeo, and Francis Rubinstein. Co-simulation based building controls implementation with networked sensors and actuators. In Proceedings of the Third ACM Workshop on Embedded Sensing Systems for Energy-Efficiency in Buildings, BuildSys '11, pages 55–60, New York, NY, USA, 2011. ACM.

Appendix A

Appendix

A.1 Collaboration

Before any actual work could start, one preliminary goal was to figure out how we could make our group work together as one. Actually this challenge is even more so in this project than in a normal work situation: No organization is in order, no predefined roles, no actual project goals and the likes. This chapter will focus on these challenges and how we tried to handle them.

A.1.1 Social Context

When we discuss the *Social Context*, we discuss the direct milieu in which the person is and how different factors can influence this person. Communication is also a part of the social context, between one to many persons, in different time zones, different cultures etc.

A.1.1.1 Common ground

Our first step to connect to our student colleagues from Kenya was to introduce ourselves via an e-mail and just shortly highlight some common information about each person from ITU, like stating name, age etc. This method is known as creating “common ground”, as introduced by Olson and Olson [29]. The term to create common ground “refers to that knowledge that the participants have in common, and they are aware that they have it in common” [29]. Common ground is not only established through simple general knowledge about each participant.

It is also created through a person's behaviour and appearance through meetings and conversations. This is often created between persons with the same temper, sense of humor and the likes. We tried to use this method as a way of getting to know our team members, to create a level of understanding and finally to create a stepping stone from which the project could evolve from. The method of creating common ground is an early way of creating a feeling of a unity, and getting to know each other.

A.1.1.2 Trust and First Impression

This initial contact was quite frustrating because of the fact that it was difficult to get a reply from some of the group members in Kenya. Only two members responded on our first e-mail, after some time, while it was rather hard to get in touch with the others. This leads directly to two different considerations in our group work: "Trust" and "First impressions matters" as introduced by Jarvenpaa et al. [19]. Trust in group work is, as the term might indicate, a value of how much the different team members trust in each other. How much does one believe that the other team members will deliver their part of the necessary work? How much does one believe that a mail will be answered? How well does the team work together? The trust between the two subgroups, Denmark and Kenya, was relatively low because of the amount -and lack of- replies and general communication. At the time being we only expect one member from Kenya to be online during our team sessions but at the same time we expect everybody from the ITU group to be online at every session. Trust is an important factor in a group setting; it is a foundation and a crucial part to solve. The group work process is highly related to how well the group function. This is also directly coherent from the first impressions that we received from the group from Kenya, with the lack of communication and willingness to participate in the project.

A.1.1.3 Collaboration Readiness

The literature for these challenges seems to agree that these sort of problems generally arise from two different topics. One being "Collaboration- and Technology Readiness" and the other "Continuities/Discontinuities". The latter part will be discussed in the subsequent section, Groupware Technologies (See section A.1.3).

Collaboration readiness is a potential show stopper for the team work, if a given member is not ready to collaborate. This could be caused by having conflicts in interest, e.g. one is about to overtake another persons job or the likes. This could cause that the person, who is about to lose his job, not to be ready to collaborate

in a productive manner. We have tried to identify these issues towards our fellow group members and we cannot find anything that should indicate that they would not be willing to collaborate. They should be just as interested in delivering a good product.

One thing that could cause their lack of interaction in our e-mail correspondences and Skype meetings are the technology readiness. We know that it is a challenge for some of the group members to get internet access. A recent study shows that only 36,3 % have access to internet in Kenya [28], which compared to Denmark's 86 % in 2010 [9], indicates that it is not as easy to get internet access in Kenya as we are used to in Denmark. And we know by fact that a few of them do not have a stable internet connection in their home. Our approach to solve this issue was to have a meeting each Tuesday at 10:00. This way we know that they should have access to their University, which most of the time has an internet connection they could use. We know they should have time for this meeting, thus it is planned as a course on their schedule.

A.1.1.4 Ethnocentrism

A potential threat to the well-being and harmony of the group is known as Ethnocentrism [8]. Ethnocentrism is a state of one subgroup, where the members sees that one group as the centre of everything, and every other group will be valued and ranked from this. A subgroup is a group inside a group, some of the group members are part of - but not the whole group is part of. This way it could create some sense of "Us versus Them", which is something you definitely want to avoid. One way to avoid this is to create multiple subgroups inside the group. One subgroup could be those who prefers to work with Java and back-end programming, while others prefer jQuery and front-end development. This could be subgroups across the different countries. If you are part of multiple subgroups the feeling of being part of just one group will dissolve, which should result in a more harmonious group. This way of creating subgroups would be something that you did early on during the initial communication, and something we tried to solve by writing small parts about each other member from Denmark. Unfortunately, we did not receive any feedback from Kenya. This has probably strengthened the feeling of us vs. them, because we do not know much about them. We definitely have a feeling that our group is the centre right now due to the fact that it is only the Danish group that develops and contributes to the project.

A.1.1.5 Coupling of work

Coupling of work relates to the state of the current tasks and how loosely or tightly coupled they are. A completely loosely coupled work is one you can perform without the interaction and feedback from other persons. A tightly coupled work task is, on the other hand, one you only can perform with other members of the group participating. Our project has evolved from a very tightly coupled project to a more loosely coupled. This was done on purpose. In the beginning of the project everybody had to be at the meetings because of the fact that we had to define which way to move the project. The development life cycle was quite rigid and strict. After the initial phase, where we decided on the platform, chose an architecture etc., more and more tasks became slightly more loosely coupled one. This means that one could start to work on his part of the project without any direct interaction with other team members. This would also allow such a highly distributed team as ours to collaborate in an efficient way. It would just be too inefficient if everybody had to be together at the same time and place every time anything regarding the project should happen. As of right now we communicate through different Groupware Tools (see A.1.3) and only meet through one weekly meeting.

A.1.2 Collaborative work

Collaborative work across cultures is a challenge. In our case we had to work together 9 people. 4 being from Kenya, a culture and country that we before entering this project, knew little about. As mentioned earlier the social context and the process of creating “common ground” with the collaborators is of high importance. The goal was to create a fundamental shared understanding of the task and build up the motivation and the much needed trust for the collaboration to succeed. Cooperative work is defined by Schmidt as “People engage in cooperative work when they are mutually dependent in their work and therefore are required to cooperate in order to get the work done” [34].

A.1.2.1 Articulation work

Articulation work is the extra activities required for collaboration [34]. The extra work is the essence of everything that is needed in order to fulfil the task, e.g. coordination of tasks. Articulation work is about who does what, when and where. There are mechanisms of interaction that supports the process -the life cycle of the project- when articulation work cannot be handled through every day social interaction. These mechanisms are for instance: Organizational structures (formalinformal), plans, schedules and conceptual schemes. Our strategy for the

articulation work was to define processes and choose the groupware technologies that supported our cause best possible.

A.1.2.2 Coordination

The more spread out between people and artefacts a task is the more the reach is increased. Increased “reach” of a task changes the coordination. Segregation, which is a method to separate a complex tasks to smaller subtasks, is a suitable strategy when a complex task is at hand. By dividing the complex task into smaller tasks you can often solve them concurrently and thereby complete the larger complex task faster. It also allows for more specialised teams to investigate a task at a more detailed level. Therefore we identified a given number of subtasks, by segregating the larger task. This could be to identify tasks such as “Create DataAccessLayer”, instead of an overall task like “Create Backend”. This is one of the ways of how we segregated the tasks within our project. Afterwards each group member were assigned to these tasks, who would work closely together, until the given subtask was completed.

To handle these tasks we created a project plan with deadlines and milestones. Moreover we agreed on having a status meeting every week, where we would discuss progress, issues, ideas etc. Arranging a meeting where all are able to attend is not always easy. We did manage to meet on Skype at times which took into consideration both the differences in time zones (temporal discontinuities, see A.1.4.1) and the fact that people had entirely different classes and work schedules.

With computer supported cooperative work (CSCW), it is impossible to anticipate every contingency. There will always be exception handling. The core challenges and dimensions of cooperative work includes articulation work, adaptation of technologies and awareness. The lack of trust and awareness when you never meet face to face with your collaborators is problematic and requires methods and training when using communication tools. We primarily used Skype to communicate with and made sure to document changes and commits to the code base very detailed. Individuals working together need to be able to gain some level of shared knowledge about each other’s activities.

A.1.3 Groupware Technologies

One way to describe collaborative software, is that it is a software designed to help people achieve a common work task within a group. One of the earliest definitions of collaborative software is found in the research paper “Rhythms, Boundaries, and Containers” by Peter and Trudy Johnson-Lenz. They describe the software as:

“intentional group processes plus software to support them.” [20]. The definition was first stated in 1974. Many things have changed, especially in the IT business, but the way a group collaborates and the group dynamics has not. Some general terms are still vital to discuss. We find the most importantly used terms in our project to be critical mass, adaptation and adoption process. These three terms and a general overview of the used tools will be explained in the sections below.

A.1.3.1 Tools

Generally we used four different tools to communicate our teamwork and progress; one synchronous and three asynchronous. We used Skype for every team meeting and for general communication between the members of the group. Skype is synchronous communication as you will normally get an instant reply while using voice chat. We furthermore used three asynchronous communication methods, that is GitHub, Email and Google Drive. GitHub was used to share the actual projects current status by communicating via tickets what needs to be done. Email was used for communicating general group announcements, that we want to make sure that all of the members from the group receive. Lastly we used Google Drive to share important documents.

A.1.3.2 Critical mass

Critical mass [14] is a sociodynamic that describes the necessary amount of people needed in order for a product to be a success. We have used this quite directly in our usage of groupware technologies, such as Skype and Google Drive. After selecting which tools to use to communicate between the different members, we have made a great deal out of actually using the tools. This has created the critical mass, for each tool, and every member knows that they can reach each other through these, and thus made it necessary for one to use.

A.1.3.3 Adaptation & Adaption Process

Adaption and the adaption process, by Tyre and Orlikowski, [37] are two closely related topics that we briefly touched through our project. These terms relates to how well an organization adapts a new technology and how they do it. Some of the problems with new technologies are that human beings tend to have routines that are not that easy to steer away from. This results in a thought of “Why should I use this new technology, when it works perfect the way I used to do it”. One of the ways we tried to solve these problems was that we chose our tools as a group, and not by enforcing it. This way the majority chose their favourite

tool. The adaption process describes how an organization adapts a new tool: The first couple of weeks are extremely important for a successful adaption. If the adaption does not succeed through these weeks it is very likely that the tool will fade out and never be used again. Our approach to this challenge has been that we pushed it to the people who did not know it, and tried to help them with the setup etc. An example of this could be the installation of Eclipse. One of the members from Kenya had some problems, and we all stayed online throughout the process and helped him install it. This way he could push it to his fellow team members in Kenya.

A.1.4 Virtual Project Management

Virtual Project management is about handling the entire project online, in a virtual and decentralized manner, where the users can access the project from any place and contribute seamlessly in developing the solution. We coordinated activities through GitHub: handling the code base, reporting and updating issues as we moved forward. The mode of interaction has mostly been through chat and voice via Skype, and not through face to face meetings. This presented various discontinuities.

A.1.4.1 Discontinuities

Watson-Manheim et al. [38] examine virtuality in terms of boundaries and discontinuities. They define discontinuities as “a break or gap in the work context”, or a “lack of continuity”. They proposed the concept of discontinuities as a general notion to permit a more comprehensive understanding of the many ways in which virtuality can be perceived. Distance is the most obvious boundary that is encountered in virtual work but it is clear that there are more boundaries such as time, organization, and nationality, which are not usually present in more conventional work settings to the same extent. It is only when those working in virtual settings perceive a boundary to be a discontinuity that it hinders work processes.

General properties of discontinuities are that they can emerge and change over time as people adapt in the teams. Discontinuities may only affect parts of the work. The typical discontinuities are temporal (working across time zones), geographic work location, work group membership (e.g. who you work with), organizational affiliation and cultural backgrounds. However discontinuities can also be expertise related (novice versus experts), historical (different version of a product), different professions (e.g. developers and researchers) or different technologies.

In our group we experienced how the different cultural backgrounds became a discontinuity. After a couple of scheduled meetings it was clear to us that the respect towards group agreements were not as important in Kenya as in Denmark. This became especially clear in regards to scheduled meetings, where the Danes often would only be a couple of minutes late, if any, which was not the case with the Kenyan members. This could be seen as a general cultural difference, or just the few members different respect of these meetings.

Another cultural difference we encountered was in regards to their view on us as Europeans. When asking Kenyans for feedback or criticism they are often fine with how it is. This might be because they do not want to cause problems and do not want to appear impolite. This could also be a part of their culture as a way to show respect.

The concept of teams varies across cultures and organizations, and how teams are perceived will differ based on the organizational and national cultural attributes of its members, according to Gibson et al [11]. It is also critical to note that individuals from different national cultures vary in terms of their group behaviours and communication styles, as Gudykunst specifies it [15].

A.1.4.2 Continuities

Continuities are the opposite of discontinuities. Continuities are the stable factors in the collaboration that the participants are aware of and consciously act upon, or they may be implicit and unrecognised [38]. Often continuities can be described as strategies or factors to overcome discontinuities.

A.2 Requirements

A.3 Tests

A.4 Project Plan

A.5 Project Plan

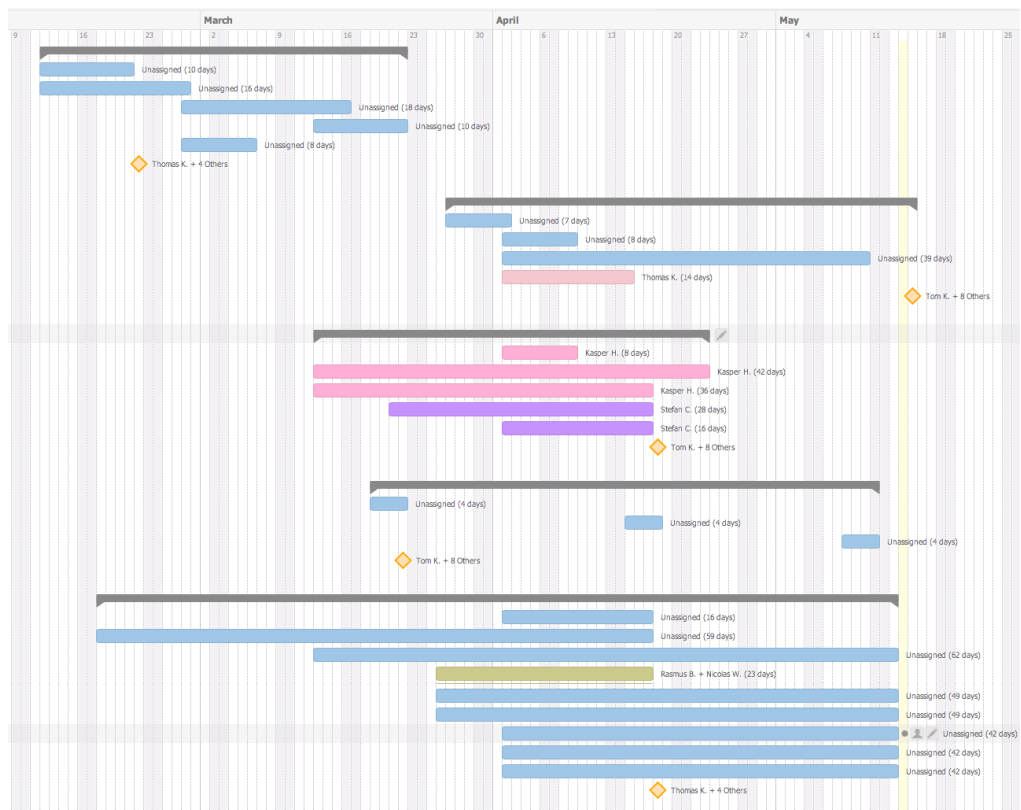


Figure A.1: Gantt Chart of the Project Plan